

recon-agent • evaluating AI agents for financial research


It tells you whether an agent change is safe to ship, what a useful answer costs, and whether each claim traces to data the agent really retrieved.

An agent answers analyst questions with tools over SEC EDGAR data. The harness scores the answer and the path taken: the harness is the product, the agent is what it measures. "Reconciliation" here means checking each claim against tool output really returned in that run, and comparing a candidate run with a committed baseline. It is not ledger or bank reconciliation.

BUSINESS PROBLEM


Analysts want agents that answer questions from filings: revenue, margins, segment figures, risk-factor text. A wrong or unsupported number in research is costly, and a new model, prompt or tool can lower accuracy or raise cost without anyone noticing. recon-agent runs each candidate on fixed benchmark questions and checks it against the baseline before it ships, prices each useful answer, and traces claims to source data. On the current baseline a case scoring at least 0.5 costs about €0.20, judge cost included.

SYSTEM ARCHITECTURE




AGENT SETUP AND VERIFICATION

MULTI MODE



VERIFICATION CHAIN



Guardrails. Tools return one of five statuses and never raise. Runs have budgets for tool calls, tokens and time. The one write path needs confirmation, and prompt-injection tests check tool data can't trigger it. Eval runs are capped at €1 by default.

CURRENT RESULTS

Runtime, mode	Cases	Task completion	Answer score	Cited rows returned	Key claims cite a returned row	Total cost
Agent SDK, single	7	1.000	0.772	1.000	0.786	€1.42

The committed baseline: 7 cases that need filing text, judged on Haiku 4.5. Task completion means a non-empty answer without a runtime error, not a correct one. Multi mode and LangGraph are not yet measured on the current rubric.

KNOWN LIMITS

- The citation check confirms a cited row was returned, not that it supports the claim. Case 731403050270 scored 0.944 with grounding rated 1.0, yet none of its key claims cites a returned row and its filing-text faithfulness is 0.0.
- The baseline covers 7 cases. The gate treats a change of up to 0.10 in answer score, or one case in completion, as run-to-run noise, so smaller effects can't be confirmed.
- LangGraph has not been run: it needs a metered API key, and the project runs on a Claude subscription only.

Source: github.com/jellewillekes/recon-agent, README, docs/architecture.md and evals/baseline.json (run eval-20261007T175852Z, rubric version 4).